

hypothesis of Thm X holds.

EEG106 LECTURE 23 Date: 21/7 Apr 2024

Classical RL \rightarrow Q-learning, SARSA
Discounted Cost Setting (S, A, P, C, α)

Assumption: Underlying MDP is unichain. For average cost MDP, so such assumption was made when the horizon was assumed to be made.

In the RL setting, we require a connectivity assumption or mistakes are compulsory and sub-optimal actions have to be taken.

Flavours of RL:
1) On-policy learning
2) Off-policy learning

On-policy \rightarrow The learning algorithm is in control and decides which action to take.

Off-policy \rightarrow The entity which decides what action to take is different than the learning agent. We can still learn from the feedback signals received by someone else.

Flavours of RL algorithms:
1) Model-based
2) Model-free

Model-based \rightarrow Algorithm is learning a "model" and hence learning (P, C)

Model-free \rightarrow Not learning P, C but directly trying to learn optimal policy. learn value functions directly.

Note that we assume both S and A to be finite. Not doing stuff such as function approximation.

Lecture 23

1 Simple model-based on-policy approach

Idea: Play random actions until all (s,a) pairs are seen
Then, Construct $\hat{P}_{ij}(a), \hat{C}(i,a)$

$$\downarrow$$
$$\frac{\# \text{ played } a \text{ and landed in } j \text{ from } i}{\# \text{ played } a \text{ from } i}$$

For the policy, play ϵ -greedy policy. i.e. w.p. ϵ_t play randomly
& w.p. $1-\epsilon_t$ play optimal action for $(S,A,\hat{P},\hat{C},\alpha)$

ϵ_t is chosen to be a decreasing function of time.

Typically, $\epsilon_t \sim \frac{1}{t}$ or slower, such as $\epsilon_t \sim \frac{1}{\sqrt{t}}$. (Borel-Cantelli Lemma)

This ensures $\hat{P} \xrightarrow{\text{a.s.}} P, \hat{C} \xrightarrow{\text{a.s.}} C \leftarrow \begin{array}{l} \text{ensure infinite exploration} \\ \text{w.p. 1} \end{array}$

The memory footprint is $O(S^2A)$

E_t is chosen to be a decreasing function of time.

Typically, $E_t \sim \frac{1}{t}$ or slower, such as $E_t \sim \frac{1}{\sqrt{t}}$. (Borel-Cantelli Lemma)

This ensures $\hat{P} \xrightarrow{\text{a.s.}} P, \hat{C} \xrightarrow{\text{a.s.}} C \longleftarrow$ ensure infinite exploration w.p. 1

The memory footprint is $O(S^2A)$

↓
storing estimates of P .

[2] Q-Learning (model-free off-policy algorithm)

learns the Q-function, with a memory requirement of $O(SA)$.

$$Q^*(i, a) = C(i, a) + \alpha \sum_j P_{ij}(a) V^*(j) \\ = C(i, a) + \alpha \sum_j P_{ij}(a) \left[\min_{a'} Q^*(j, a') \right]$$

This is the Bellman equation in terms of the Q function

$$\hat{Q}(x_t, A_t) \leftarrow (1 - \gamma_t) \hat{Q}(x_t, A_t) + \gamma_t [C_t + \min_{a'} \hat{Q}(x_{t+1}, a')]$$

γ_t is the step-size parameter, gives memory of the earlier event

γ_t is also chosen as a decreasing function of time

Note that $\min_{a'} \hat{Q}(x_{t+1}, a')$ is an estimate of $V^*(x_{t+1})$.

Need: $\sum \gamma_t = \infty, \sum \gamma_t^2 < \infty$

Q: What do we need for \hat{Q} to converge to Q ?
 → In the long run, every (s, a) -pair should be visited almost surely
 Thus, we cannot have a stationary policy which drives Q -learning

"Tracks ODE" : $\dot{Q} = TQ - Q$, where T is an operator

$$T_u(i, a) = C(i, a) + \alpha \sum_j P_{ij}(a) \left[\min_{a'} u(j, a') \right]$$

 ↓
 a candidate function

Note that the Bellman equation can be expressed as $TQ = Q$
 The memory footprint now is $O(s, a)$.

3 SARSA (State, Action, Reward, State, Action)
 Model-free, on-policy algorithm

$A_t = \begin{cases} \text{Play random action w.p. } \epsilon_t \\ \arg \min_a \hat{Q}(x_t, a) \text{ w.p. } 1 - \epsilon_t \end{cases}$

$$\hat{Q}(x_t, A_t) \leftarrow (1 - \gamma_t) \hat{Q}(x_t, A_t) + \gamma_t [C_t + \alpha \hat{Q}(x_{t+1}, A_{t+1})]$$

There is no minimization term in SARSA. Note that the next step we WILL perform minimization w.p. $1 - \epsilon_t$ action selection

$$A_t = \begin{cases} \text{Play random action w.p. } \epsilon \\ \arg \min_a \hat{Q}(X_t, a) \text{ w.p. } 1 - \epsilon \end{cases}$$

$$\hat{Q}(X_t, A_t) \leftarrow (1 - \gamma_t) \hat{Q}(X_t, A_t) + \gamma_t [C_t + \alpha \hat{Q}(X_{t+1}, A_{t+1})]$$

There is no minimization term in SARSA. Note that the next step we WILL perform minimization w.p. $1 - \epsilon_t$

Note that for Q -learning, we can use any action selection algorithm as long as it performs infinite exploration.

Typically choose $\epsilon_t \sim \frac{1}{t}$ & $\gamma_t \sim \frac{1}{t}$
 $\Rightarrow \hat{Q} \xrightarrow{\text{a.s.}} Q$

Note that these algorithms are notoriously slow. This is because only entry gets filled in one time-step

Function approximation: expresses value function as a linear combination of some basis functions.

4 Actor-Critic algorithm

GPI \rightarrow Generalized Policy Iteration

Original policy iteration searches for greedy improvements on the policy

Actor: Policy improvement (slower timescale)

Critic: Policy evaluation (faster timescale)

You have two timescales, one with a small & one with a large step size
In practice, use step size γ_t for actor & β_t for critic.

$$\beta_t = o(\gamma_t) \quad \gamma_t = o(\beta_t)$$

↓
Little \approx Oh

SARSA

If we do, off-policy SARSA, under some stationary policy π , we will learn the value function corresponding to the policy $\pi \rightarrow Q^\pi$ and V^π .
All of these algorithms can be generalized for a time-average MDP setting.